

Compressed Representations of Set Collections

Laurea Magistrale in Informatica

Luca Lombardo

29 May 2026

Dipartimento di Informatica

Università di Pisa

Set collections

Definition (Set collection)

A set collection over a finite, totally ordered universe U is a family

$$\mathcal{S} = \{S_1, \dots, S_m\}, \quad S_i \subseteq U.$$

We write $u = |U|$ and $n = \sum_{i=1}^m |S_i|$, and assume $U = [1..u]$.

The encoded object is the whole collection, but each query selects one set S_i and asks for ordered-set information inside it.

- **member**(S_i, x): whether $x \in S_i$.
- **access**(S_i, j): the j -th element of S_i .
- **rank**(S_i, x): the number of elements of S_i at most x .
- **predecessor**(S_i, x): the largest element of S_i at most x .
- **successor**(S_i, x): the smallest element of S_i at least x .

Counting is not querying

A representation identifies one object among many possible objects. A bit is redundant when changing it does not change which object is identified.

Proposition (Counting lower bound)

Let $\mathcal{F}$ be a finite class of objects. If every object in $\mathcal{F}$ is represented by a binary codeword of length at most L , then

$$2^L \geq |\mathcal{F}| \quad \text{and therefore} \quad L \geq \lceil \log |\mathcal{F}| \rceil.$$

The bound is only a **space target**. A shortest description may identify the object, but it may give no direct way to test membership, count elements, or select the next one.

We seek representations close to the counting cost that still answer queries.

Independent encoding

Per set. If relations inside $\mathcal{S}$ are ignored, each set is chosen as a subset of the full universe.

$$S_i \subseteq U, \quad |S_i| \text{ fixed} \implies \binom{u}{|S_i|} \text{ choices.}$$

Whole collection. Describing the sets independently adds the per-set costs, giving the baseline space for the collection.

$$H_{\text{wc}}(\mathcal{S}) = \sum_{i=1}^m \lg \binom{u}{|S_i|}$$

Independent encoding pays for each set in isolation, as if the other sets were not present. Any smaller description must use redundancy across the collection.

Containment references

A first way to use redundancy is to describe a set inside a containing reference.

Definition (Containment step)

$$Q \subseteq P \implies \text{cost} = \lg \binom{|P|}{|Q|} \leq \lg \binom{u}{|Q|}.$$

Alanko et al. (2025) build a **containment hierarchy** where each set is represented inside a smallest proper superset, or inside U if no such set exists.

$$U \rightarrow \dots \rightarrow P \rightarrow Q, \quad Q = P \setminus D.$$

This is query-friendly because every parent-to-child step is a deletion. It becomes compressive when the collection contains a **deep nesting** structure.

Symmetric differences

If neither $A \subseteq B$ nor $B \subseteq A$, containment cannot relate the two sets directly, even when they almost coincide.


Definition (Symmetric difference)

$$A \triangle B = (A \setminus B) \cup (B \setminus A), \quad |A \triangle B| = |A \setminus B| + |B \setminus A|$$

If S is described from a chosen reference R , only the changed elements are stored as signed edits.

$$+x \text{ for } x \in S \setminus R, \quad -y \text{ for } y \in R \setminus S$$

Each reference costs $|S \triangle R|$. Gagie, He, and Navarro (2026) choose the references for the whole collection, then store the resulting signed edits on two rooted trees, called **indel trees**.



Choosing references

To query a set through a path, we first choose a reference for it. That choice fixes what the path records: deletions for containment, signed edits for symmetric difference.


Structure	Reference domain	Stored differences
Containment	containing set $P \supseteq S$	deletions $P \setminus S$
Symmetric difference	nearby existing set R	signed edits $S \triangle R$

available references $\rightarrow$ stored differences $\rightarrow$ queryable paths

Keep deletion paths, but allow new containing references.

If two sets overlap but neither contains the other, the pair itself gives no containment reference. The smallest containing reference we can add is their union.


Union as reference



Encode A and B from M

First bitvector: marks $A \cap B$ inside M

Second bitvector: on $M \setminus (A \cap B)$, 1 marks $A \setminus B$




Once the union M is stored, the bitvectors encode how M splits into the three regions of (A, B) . This local split costs $w(A, B)$.

$$w(A, B) = \lg \binom{|M|}{k, l, r} + c(|M|, k) \quad \text{where} \quad k = |A \cap B|, \quad l = |A \setminus B|, \quad r = |B \setminus A|.$$

Set-Union Matching (SUM)


The construction builds a **bottom-up hierarchy** from the input sets: each round scores every pair in the current family and replaces the selected pairs by their unions.




We need a **score** for each candidate merge, so that each round can choose as many pairs as possible without using the same current set twice and with minimum total cost.

Merge score

Old description



New description





$$\Delta\Phi(A, B) = \underbrace{\lg\left(\frac{u}{|M|}\right)}_{\text{store the union}} + \underbrace{w(A, B)}_{\text{split into } A \text{ and } B} - \underbrace{\lg\left(\frac{u}{|A|}\right)}_{\text{old support of } A} - \underbrace{\lg\left(\frac{u}{|B|}\right)}_{\text{old support of } B}.$$

$\Delta\Phi < 0$: **saves bits**

$\Delta\Phi = 0$: same counted cost

$\Delta\Phi > 0$: **not locally profitable**


Greedy matching



$$\Phi(F^*) = \min_t \Phi(F_t) \leq \Phi(F_0) = H_{wc}(\mathcal{S}).$$

Because F_0 has no merges, the independent encoding is always among the recorded candidates.

Recorded forest



Deletion paths

Consider the path from a component root R to the selected input set S , with intermediate sets $P_1, \dots, P_{t-1}$. Since each parent is a union, this path is a **containment chain**.

$$R = P_0 \supseteq P_1 \supseteq \dots \supseteq P_t = S$$

For each edge, **store the elements deleted** at that step.

$$D_i = P_{i-1} \setminus P_i$$


The selected input set is recovered from the component root by deleting the elements encountered on the path:

$$S = R \setminus \bigcup_{i=1}^t D_i$$

An element deleted on the path cannot reappear below the edge where it is removed.

Survivor queries

ranks in root
 R_i



path to S_i
deleted local ranks



kept ranks
define S_i



$$D(v_i) = \{2, 4, 7, 8\} \text{ and } S_i = \{x_j : j \in [1..9] \setminus D(v_i)\} = \{x_1, x_3, x_5, x_6, x_9\}$$

member(S_i, x)

$j = \text{rank}_{R_i}(x)$
true iff rank j survives

rank(S_i, x)

$j = \text{rank}_{R_i}(x)$
 $j - |D(v_i) \cap [1..j]|$

access(S_i, q)

select the q -th survivor rank j
return $\text{select}_{R_i}(j)$

Query times

After the directory finds (R_i, v_i) , the root dictionary translates universe values to ranks in R_i . The remaining cost is a path query over the local deletion alphabet $[1..|R_i|]$.

Theorem (SUM query times)

For nonconstant $|R_i|$, the representation supports:

Operation	Question on the path	Time
$\text{member}(S_i, x)$	<i>was the rank of x deleted?</i>	$O(\lg \lg_\omega R_i)$
$\text{rank}(S_i, x)$	<i>how many deleted ranks lie before x?</i>	$O(\lg_\omega R_i)$
$\text{access}(S_i, j)$	<i>which rank is the j-th survivor?</i>	$O\left(\frac{\lg R_i }{\lg \lg R_i }\right)$
<i>predecessor, successor</i>	<i>rank followed by access</i>	<i>same as access</i>

These bounds depend on the **component root size** $|R_i|$, not directly on the universe size u . If $|R_i| \ll u$, the logarithmic factors are smaller.

Queryable storage

Theorem (Representation space)

Each component root R is stored directly as a rank/select dictionary over U .

$$\sum_{R \text{ root}} \left(\lg \binom{u}{|R|} + o(u) + O(\lg \lg u) \right) \text{ bits}$$

The overlay stores the input directory and the deletion labels generated by the internal occurrences. Here m is the number of input sets.

$$O \left(m + \sum_{c \text{ internal}} |P_{\text{left}(c)} \triangle P_{\text{right}(c)}| \right) \text{ words}$$

These data support the five selected-set queries.

Compressed Representations of Set Collections

Thank you for listening!
Any question?

Complement selection

Path selection returns labels present on the path, hence deleted ranks. Access must select a **survivor**, so it asks for a rank absent from that path.

Definition (Complement path selection)

Let $D_{T(v)} \subseteq [1..\sigma]$ be the deletion labels on the root-to- v path. Then $\text{path_compselect}(v, j)$ returns the j -th smallest label in $[1..\sigma] \setminus D_{T(v)}$.

In an interval J , path counting gives the number $C_{v(J)}$ of deleted labels, so the survivor count is

$$|J| - C_{v(J)}.$$

The same alphabet descent supports complement selection in $O\left(\frac{\lg \sigma}{\lg \lg \sigma}\right)$ time for $\sigma \geq 2$, and in constant time for $\sigma \leq 1$. For SUM, $\sigma = |R_i|$, and access returns

$$\text{select}_{R_i}(\text{path_compselect}(v_i, j))$$

Membership patterns

After building a queryable hierarchy, we can ask what any representation is trying to approach. **Ignore references** for a moment and describe the collection element by element.

Definition (Membership pattern)

$$\sigma(x) = \{i \in [m] : x \in S_i\}.$$

Equal patterns identify elements that no input set can distinguish. For each pattern $\tau \subseteq [m]$, collect all such elements in one region.

$$A_\tau = \{x \in U : \sigma(x) = \tau\}.$$

The regions A_τ partition U . Their sizes count how many times the collection repeats each membership pattern, independently of any chosen hierarchy.

Atom bound

Once the pattern regions and their sizes are fixed, the remaining choice is how the labelled elements of U are assigned to those regions.

Definition (Realised patterns and atom sizes)

$$R(\mathcal{S}) = \{\tau \subseteq [m] : A_\tau \neq \emptyset\}, \quad c_\tau = |A_\tau| \quad \text{for } \tau \in R(\mathcal{S}).$$

The assignment must respect that exactly c_τ elements receive each realised pattern τ .

Proposition (Atom bound)

With this side information, every lossless representation needs at least

$$L^*(\mathcal{S}) = \lg \binom{u}{(c_\tau)_{\tau \in R(\mathcal{S})}} \text{ bits.}$$

Global code

The lower bound is attainable as a global enumerative code.

Proposition (Achievability of L^*)

Given $R(\mathcal{S})$ and $(c_\tau)_{\tau \in R(\mathcal{S})}$, the collection can be encoded in exactly $\lceil L^(\mathcal{S}) \rceil$ bits.*

Order $U = \{x_1, \dots, x_u\}$. The encoder maps the collection to the word of membership patterns along this order. Its alphabet is $R(\mathcal{S})$, and c_τ fixes how many times each symbol τ appears.

$$\pi(\mathcal{S}) = (\sigma(x_1), \dots, \sigma(x_u)) \in R(\mathcal{S})^u \quad |\{j : \sigma(x_j) = \tau\}| = c_\tau$$

The code is optimal because it identifies this word among all words with the same counts. It remains global, while SUM is queryable because it exposes local root-to-leaf paths.