



Efficient Succinct Data Structures on Directed Acyclic Graphs

Tesi Triennale in Matematica

Luca Lombardo

9 Maggio 2025





Why Succinct Data Structures?

Massive Data and Structure Overhead

The Challenge: Massive Data & Query Needs

Modern datasets (Science, Web, AI...) are enormous. Complex analysis demands data in RAM, but auxiliary structures (indexes, trees) needed for queries often **occupy more space than the data itself.** \implies Fitting everything in RAM is a major bottleneck.



Why Succinct Data Structures?

Massive Data and Structure Overhead

The Challenge: Massive Data & Query Needs

Modern datasets (Science, Web, AI...) are enormous. Complex analysis demands data in RAM, but auxiliary structures (indexes, trees) needed for queries often **occupy more space than the data itself.** \implies Fitting everything in RAM is a major bottleneck.

Classic Trade-off:

- *Compression:* Minimal space, but slow/no direct queries.
- *Traditional Data Structures:* Fast queries, but large space overhead.



Why Succinct Data Structures?

Massive Data and Structure Overhead

The Challenge: Massive Data & Query Needs

Modern datasets (Science, Web, AI...) are enormous. Complex analysis demands data in RAM, but auxiliary structures (indexes, trees) needed for queries often **occupy more space than the data itself.** \implies Fitting everything in RAM is a major bottleneck.

Classic Trade-off:

- *Compression:* Minimal space, but slow/no direct queries.
- *Traditional Data Structures:* Fast queries, but large space overhead.

The Succinct Goal: Best of Both Worlds

Can we achieve **both**?

- Space **near information-theoretic minimum.**
- Efficient queries **directly** on compact data.



Shannon Entropy

Fundamental Limits of Lossless Compression

Goal: Determine the minimum average number of bits per symbol required for a lossless representation of data from a source X .



Shannon Entropy

Fundamental Limits of Lossless Compression

Goal: Determine the minimum average number of bits per symbol required for a lossless representation of data from a source X .

Definition (Shannon Entropy $H(X)$)

The average uncertainty, or information content, per symbol of source X :

$$H(X) = E_{P_X}[-\log_2 P_X(x)] = - \sum_{x \in \mathcal{X}} P_X(x) \log_2 P_X(x) \quad [\text{bits/symbol}]$$



Shannon Entropy

Fundamental Limits of Lossless Compression

Goal: Determine the minimum average number of bits per symbol required for a lossless representation of data from a source X .

Definition (Shannon Entropy $H(X)$)

The average uncertainty, or information content, per symbol of source X :

$$H(X) = E_{P_X}[-\log_2 P_X(x)] = - \sum_{x \in \mathcal{X}} P_X(x) \log_2 P_X(x) \quad [\text{bits/symbol}]$$

Source Coding Theorem (Lower Bound)

Shannon proved that $H(X)$ constitutes the **theoretical lower bound** on the average number of bits per symbol required to represent the output of source X without loss of information.



Zero-Order Empirical Entropy

A Practical Bound Based on Data

We usually don't know the true source P_X . We only have the data sequence S .



Zero-Order Empirical Entropy

A Practical Bound Based on Data

We usually don't know the true source P_X . We only have the data sequence S .

Definition (Zero-Order Empirical Entropy $\mathcal{H}_0(S)$)

Information content of sequence S based on its symbol counts (n_s):

$$\mathcal{H}_0(S) = \sum_{s \in \Sigma} \frac{n_s}{n} \log_2 \frac{n}{n_s} \quad [\text{bits/symbol}]$$

Uses observed frequencies $\frac{n_s}{n}$ instead of unknown $P_X(x)$.



Zero-Order Empirical Entropy

A Practical Bound Based on Data

We usually don't know the true source P_X . We only have the data sequence S .

Definition (Zero-Order Empirical Entropy $\mathcal{H}_0(S)$)

Information content of sequence S based on its symbol counts (n_s):

$$\mathcal{H}_0(S) = \sum_{s \in \Sigma} \frac{n_s}{n} \log_2 \frac{n}{n_s} \quad [\text{bits/symbol}]$$

Uses observed frequencies $\frac{n_s}{n}$ instead of unknown $P_X(x)$.

Relevance for Succinct Structures

$n \cdot \mathcal{H}_0(S)$ is a practical space benchmark. Succinct data structures often target space close to this value (or higher-order versions \mathcal{H}_k) for the **given sequence S**.



Bitvectors and Fundamental Queries

The Simplest Sequence

Consider the most basic sequence: a **bitvector** $B[1..n]$, a sequence of n bits from $\{0, 1\}$.

1	0	1	1	0	1	0	0	1	1	0	1	0	1	1	0	0	0	1	0
1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20



Bitvectors and Fundamental Queries

The Simplest Sequence

Consider the most basic sequence: a **bitvector** $B[1..n]$, a sequence of n bits from $\{0, 1\}$.

- $\text{rank}_b(B, i)$: How many bits b are in the prefix $B[1..i]$? (Count)
- $\text{select}_b(B, j)$: What is the position (index) of the j -th occurrence of bit b ? (Locate)

1	0	1	1	0	1	0	0	1	1	0	1	0	1	1	0	0	0	1	0
1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20

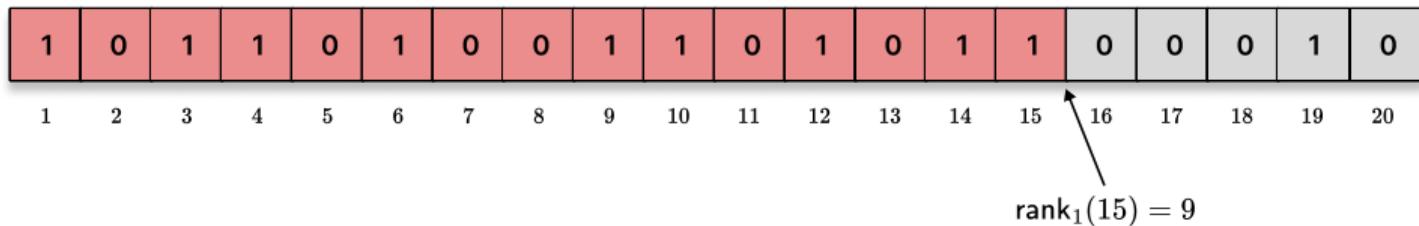


Bitvectors and Fundamental Queries

The Simplest Sequence

Consider the most basic sequence: a **bitvector** $B[1..n]$, a sequence of n bits from $\{0, 1\}$.

- $\text{rank}_b(B, i)$: How many bits b are in the prefix $B[1..i]$? (Count)
- $\text{select}_b(B, j)$: What is the position (index) of the j -th occurrence of bit b ? (Locate)



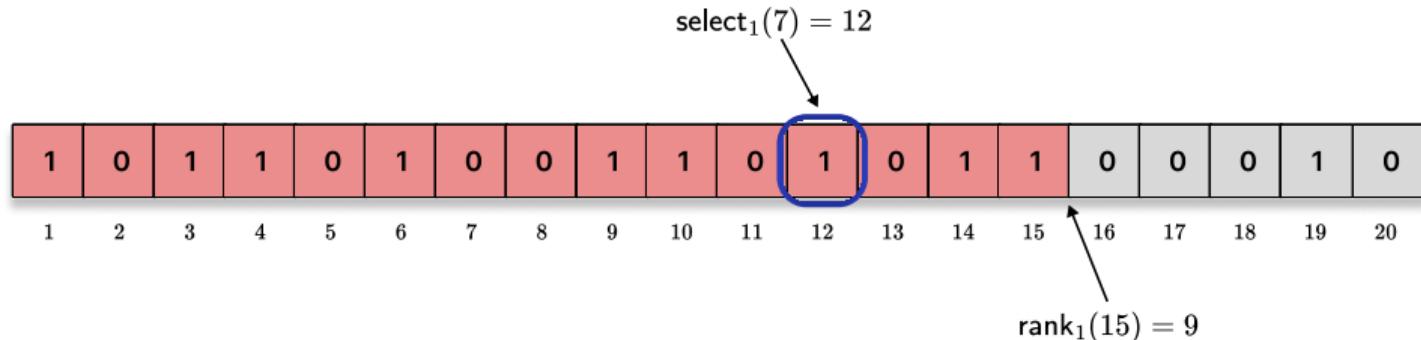


Bitvectors and Fundamental Queries

The Simplest Sequence

Consider the most basic sequence: a **bitvector** $B[1..n]$, a sequence of n bits from $\{0, 1\}$.

- $\text{rank}_b(B, i)$: How many bits b are in the prefix $B[1..i]$? (Count)
- $\text{select}_b(B, j)$: What is the position (index) of the j -th occurrence of bit b ? (Locate)



Furthermore, rank and select are inverse operations:

$$\text{rank}_b(B, \text{select}_b(B, j)) = j \quad \text{and} \quad \text{select}_b(B, \text{rank}_b(B, i)) = i.$$



RRR Structure: Entropy-Compressed Bitvectors

Achieving Space Close to Empirical Entropy

Succinct Data Structure for Bitvectors

- **Goal:** Support rank and select in $O(1)$ time.
- **Space:** Close to the information-theoretic minimum for the bitvector.



RRR Structure: Entropy-Compressed Bitvectors

Achieving Space Close to Empirical Entropy

Succinct Data Structure for Bitvectors

- **Goal:** Support rank and select in $O(1)$ time.
- **Space:** Close to the information-theoretic minimum for the bitvector.

Theorem (RRR Structure)

A bitvector $B[1..n]$ with m set bits can be represented using

$$B(n, m) + o(n) + O(\log \log n) \text{ bits},$$

where $B(n, m) = \lceil \log_2 \binom{n}{m} \rceil$, while supporting rank and select queries in $O(1)$ time.

$B(n, m) \approx n\mathcal{H}_0(B)$ is the **information-theoretic minimum space** required to store an arbitrary subset of size m from a universe of size n



Beyond Bitvectors: General Alphabets

Wavelet Trees

What about sequences $S[1..n]$ over larger alphabets $\Sigma = \{1, \dots, \sigma\}$?



Beyond Bitvectors: General Alphabets

Wavelet Trees

What about sequences $S[1..n]$ over larger alphabets $\Sigma = \{1, \dots, \sigma\}$?

abracadabra

{a, b, c, d, r}



Beyond Bitvectors: General Alphabets

Wavelet Trees

What about sequences $S[1..n]$ over larger alphabets $\Sigma = \{1, \dots, \sigma\}$?

abracadabra
00100010010

{a, b, c | d, r}

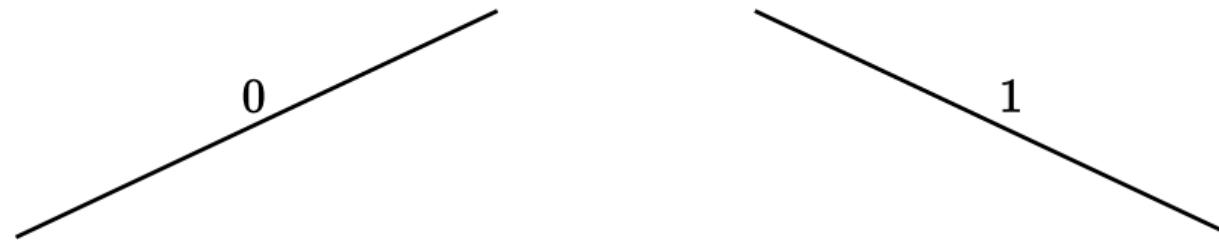


Beyond Bitvectors: General Alphabets

Wavelet Trees

What about sequences $S[1..n]$ over larger alphabets $\Sigma = \{1, \dots, \sigma\}$?

abracadabra
00100010010 $\{a, b, c \mid d, r\}$



abacaaba
00010000 $\{a, b \mid c\}$

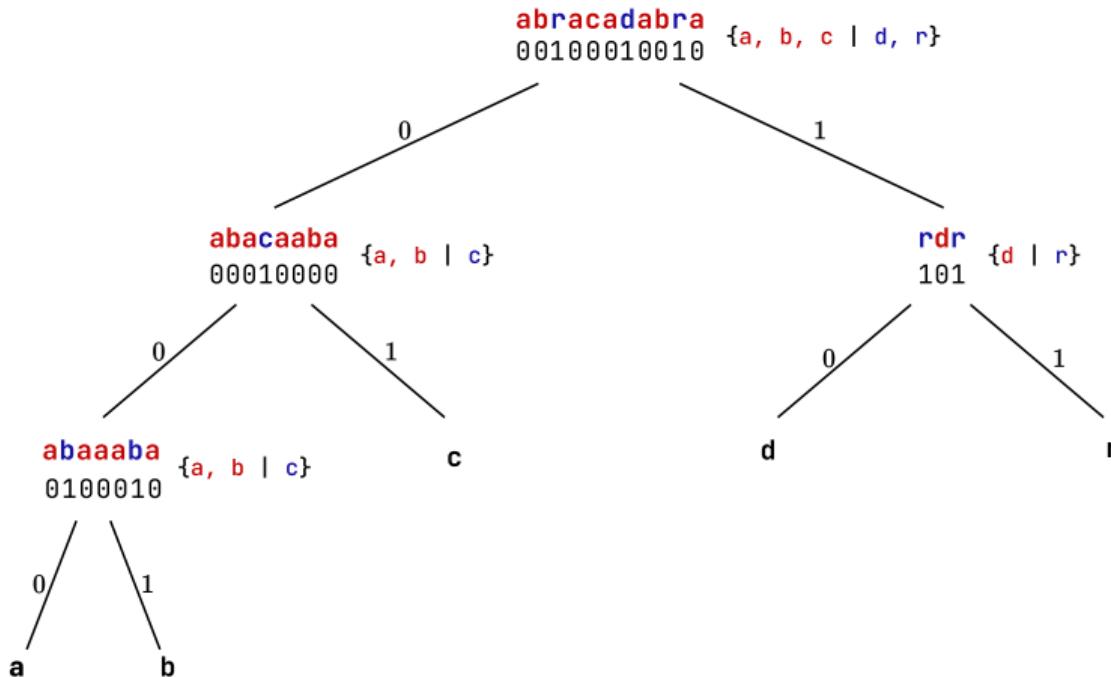
rdr
101 $\{d \mid r\}$



Beyond Bitvectors: General Alphabets

Wavelet Trees

What about sequences $S[1..n]$ over larger alphabets $\Sigma = \{1, \dots, \sigma\}$?

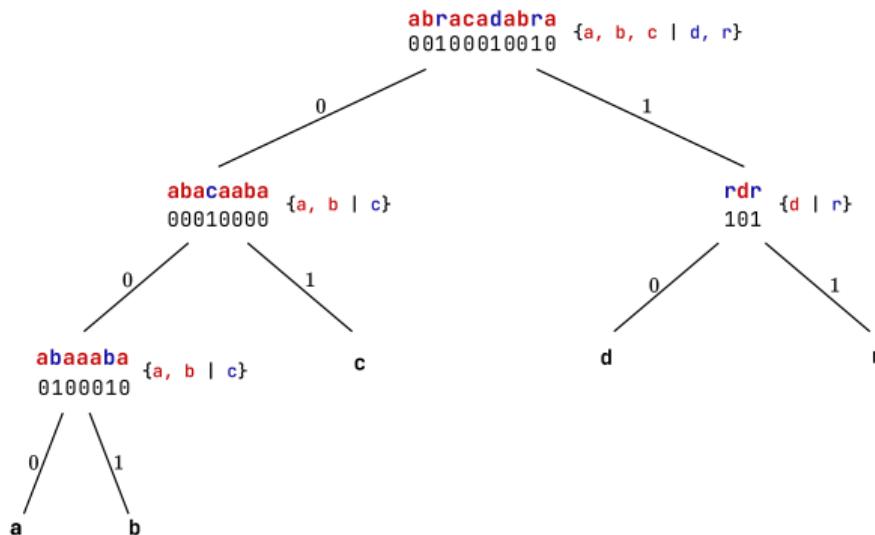




Beyond Bitvectors: General Alphabets

Wavelet Trees

What about sequences $S[1..n]$ over larger alphabets $\Sigma = \{1, \dots, \sigma\}$?



\mathcal{H}_0 -Compressed Wavelet Tree

Using RRR for bitvectors:

- **Space:** $n\mathcal{H}_0(S) + o(n \log \sigma)$ bits.
- **Query Time:** $O(\log \sigma)$ for access, rank_c , select_c .

Adapts space to the sequence's zero-order entropy.



Representing Sequence Variation: Degenerate Strings

Definitions and Core Operations

Definition and Rank & Select Adaptation

A **degenerate string** is a sequence $X = X_1X_2 \dots X_n$, where each X_i is a subset of the alphabet Σ with cardinality σ . We can define the following operations:

- $\text{subset-rank}_X(i, c)$: Counts sets X_k ($k \leq i$) where $c \in X_k$.
- $\text{subset-select}_X(j, c)$: Finds index k of the j -th set X_k where $c \in X_k$.



Representing Sequence Variation: Degenerate Strings

Definitions and Core Operations

Definition and Rank & Select Adaptation

A **degenerate string** is a sequence $X = X_1 X_2 \dots X_n$, where each X_i is a subset of the alphabet Σ with cardinality σ . We can define the following operations:

- $\text{subset-rank}_X(i, c)$: Counts sets X_k ($k \leq i$) where $c \in X_k$.
- $\text{subset-select}_X(j, c)$: Finds index k of the j -th set X_k where $c \in X_k$.

$$X = \begin{array}{cccc} \left\{ \begin{array}{c} A \\ C \\ G \end{array} \right\} & \left\{ \begin{array}{c} A \\ T \end{array} \right\} & \left\{ \begin{array}{c} C \end{array} \right\} & \left\{ \begin{array}{c} T \\ G \end{array} \right\} \\ X_1 & X_2 & X_3 & X_4 \end{array} \quad S = \begin{array}{ccccc} ACG & AT & C & TG \\ R = \begin{array}{ccccc} 100 & 10 & 1 & 10 & 1 \\ S_1 & S_2 & S_3 & S_4 \end{array} \end{array}$$



Representing Sequence Variation: Degenerate Strings

Definitions and Core Operations

Definition and Rank & Select Adaptation

A **degenerate string** is a sequence $X = X_1 X_2 \dots X_n$, where each X_i is a subset of the alphabet Σ with cardinality σ . We can define the following operations:

- $\text{subset-rank}_X(i, c)$: Counts sets X_k ($k \leq i$) where $c \in X_k$.
- $\text{subset-select}_X(j, c)$: Finds index k of the j -th set X_k where $c \in X_k$.

$$X = \begin{array}{cccc} \left\{ \begin{array}{c} A \\ C \\ G \end{array} \right\} & \left\{ \begin{array}{c} A \\ T \end{array} \right\} & \left\{ \begin{array}{c} C \end{array} \right\} & \left\{ \begin{array}{c} T \\ G \end{array} \right\} \\ X_1 & X_2 & X_3 & X_4 \end{array} \quad S = \begin{array}{ccccc} ACG & AT & C & TG \\ R = & 100 & 10 & 1 & 10 & 1 \\ & S_1 & S_2 & S_3 & S_4 \end{array}$$

To compute $\text{subset-rank}_X(i, c)$: Let p be the end position in S for prefix $X_1..X_i$, found using $\text{select}_1(R, i + 1)$. The result is $\text{rank}_c(S, p)$.



From Degenerate Strings to Weighted DAGs

A New Perspective

Recall our degenerate string X :

$$X = \begin{matrix} \left\{ \begin{matrix} A \\ C \\ G \end{matrix} \right\} & \left\{ \begin{matrix} A \\ T \end{matrix} \right\} & \left\{ \begin{matrix} T \\ C \\ A \end{matrix} \right\} & \left\{ \begin{matrix} A \\ G \end{matrix} \right\} \\ X_1 & X_2 & X_3 & X_4 \end{matrix}$$



From Degenerate Strings to Weighted DAGs

A New Perspective

Recall our degenerate string X :

$$X = \begin{matrix} \left\{ \begin{matrix} A \\ C \\ G \end{matrix} \right\} & \left\{ \begin{matrix} A \\ T \end{matrix} \right\} & \left\{ \begin{matrix} T \\ C \\ A \end{matrix} \right\} & \left\{ \begin{matrix} A \\ G \end{matrix} \right\} \\ X_1 & X_2 & X_3 & X_4 \end{matrix}$$

Key Idea

We can model queries on X by constructing a specific node-weighted Directed Acyclic Graph (DAG) G_c for a target character c .



From Degenerate Strings to Weighted DAGs

A New Perspective

Recall our degenerate string X :

$$X = \begin{matrix} \left\{ \begin{matrix} A \\ C \\ G \end{matrix} \right\} & \left\{ \begin{matrix} A \\ T \end{matrix} \right\} & \left\{ \begin{matrix} T \\ C \\ A \end{matrix} \right\} & \left\{ \begin{matrix} A \\ G \end{matrix} \right\} \\ X_1 & X_2 & X_3 & X_4 \end{matrix}$$

Key Idea

We can model queries on X by constructing a specific node-weighted Directed Acyclic Graph (DAG) G_c for a target character c .

- **Vertices V_c :** Source s + one node $v_{k,a}$ for each character $a \in X_k$ at each position k .



From Degenerate Strings to Weighted DAGs

A New Perspective

Recall our degenerate string X :

$$X = \begin{matrix} \left\{ \begin{matrix} A \\ C \\ G \end{matrix} \right\} & \left\{ \begin{matrix} A \\ T \end{matrix} \right\} & \left\{ \begin{matrix} T \\ C \\ A \end{matrix} \right\} & \left\{ \begin{matrix} A \\ G \end{matrix} \right\} \\ X_1 & X_2 & X_3 & X_4 \end{matrix}$$

Key Idea

We can model queries on X by constructing a specific node-weighted Directed Acyclic Graph (DAG) G_c for a target character c .

- **Vertices V_c :** Source s + one node $v_{k,a}$ for each character $a \in X_k$ at each position k .
- **Weights w_c :** $w_c(s) = 0$. $w_c(v_{k,a}) = 1$ if $a = c$, else 0. (Highlights occurrences of c).



From Degenerate Strings to Weighted DAGs

A New Perspective

Recall our degenerate string X :

$$X = \begin{matrix} \left\{ \begin{matrix} A \\ C \\ G \end{matrix} \right\} & \left\{ \begin{matrix} A \\ T \end{matrix} \right\} & \left\{ \begin{matrix} T \\ C \\ A \end{matrix} \right\} & \left\{ \begin{matrix} A \\ G \end{matrix} \right\} \\ X_1 & X_2 & X_3 & X_4 \end{matrix}$$

Key Idea

We can model queries on X by constructing a specific node-weighted Directed Acyclic Graph (DAG) G_c for a target character c .

- **Vertices V_c :** Source s + one node $v_{k,a}$ for each character $a \in X_k$ at each position k .
- **Weights w_c :** $w_c(s) = 0$. $w_c(v_{k,a}) = 1$ if $a = c$, else 0. (Highlights occurrences of c).
- **Edges E_c :** Connect s to $k = 1$. Connect all $v_{k,a}$ to all $v_{k+1,b}$. (Represents sequence adjacency).



Path Weight Aggregation: The \mathcal{O} -Set

Capturing All Path Weights

Given a path $P = (v_0 = s, \dots, v_k = v)$ we define $W(P) = \sum_{j=1}^k w(v_j)$

Goal

Characterize the set of **all possible distinct** cumulative path weights arriving at each node.



Path Weight Aggregation: The \mathcal{O} -Set

Capturing All Path Weights

Given a path $P = (v_0 = s, \dots, v_k = v)$ we define $W(P) = \sum_{j=1}^k w(v_j)$

Goal

Characterize the set of **all possible distinct** cumulative path weights arriving at each node.

\mathcal{O} -Set Definition (Recursive)

- **Base Case (Source):** $\mathcal{O}_s = \{0\}$
- **Recursive Step ($v \neq s$):**

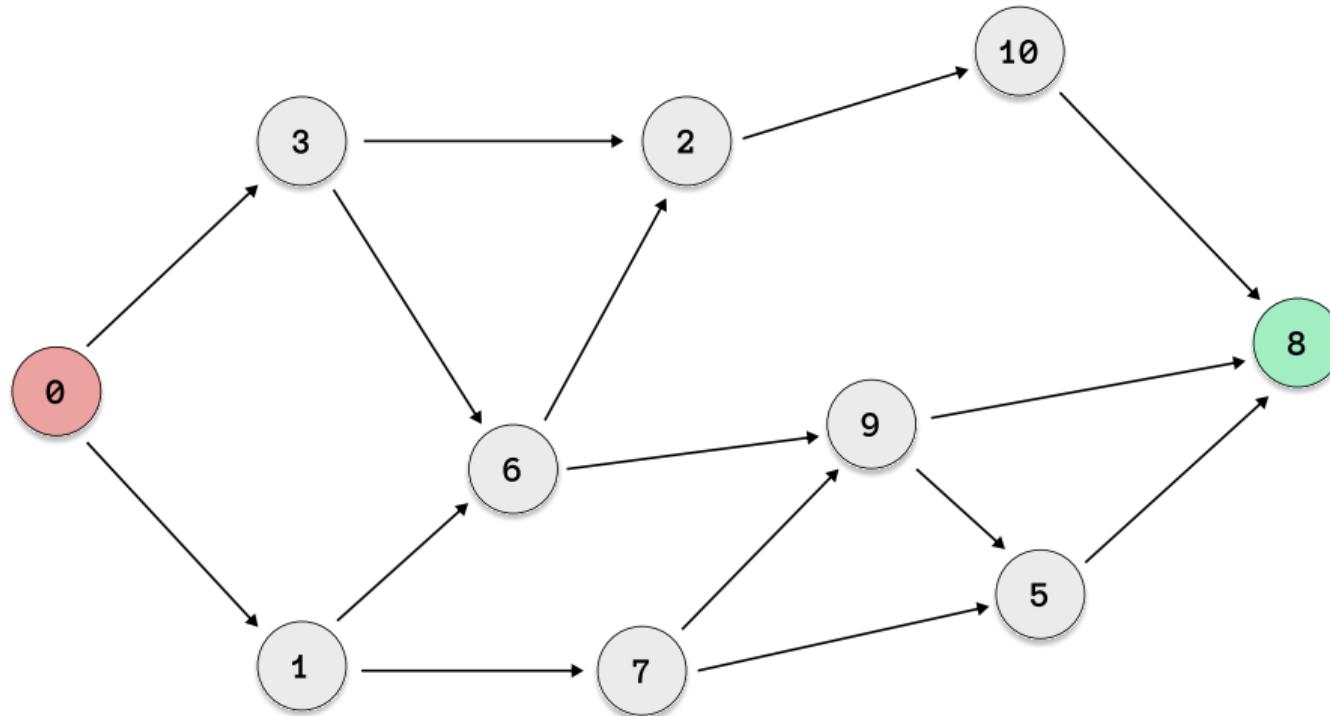
$$\mathcal{O}_v = \bigcup_{u \in \text{Pred}(v)} \{\gamma + w(v) \mid \gamma \in \mathcal{O}_u\} = \{W(P) \mid P \in \text{Path}(s, v)\}$$

Keep only **distinct** values. Store as a sorted sequence.



O-Set Construction Example

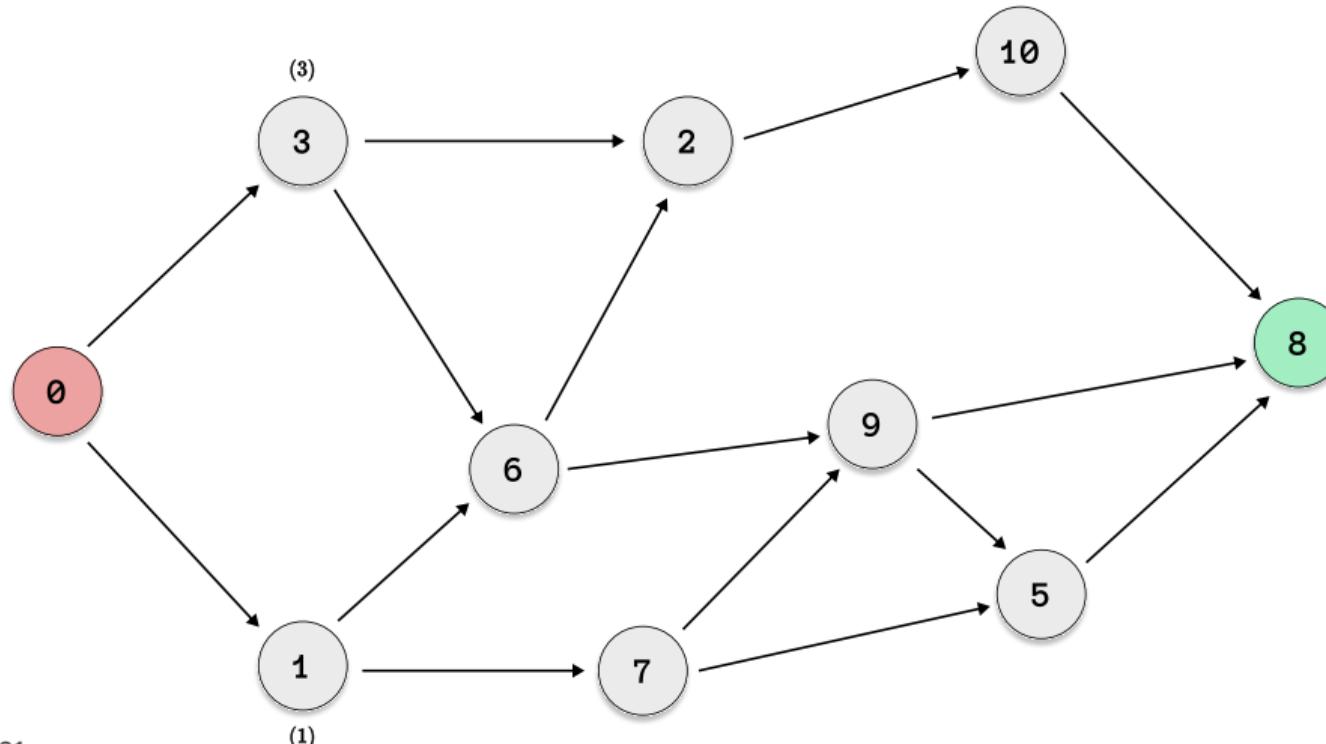
Visualizing the Recursive Definition





O-Set Construction Example

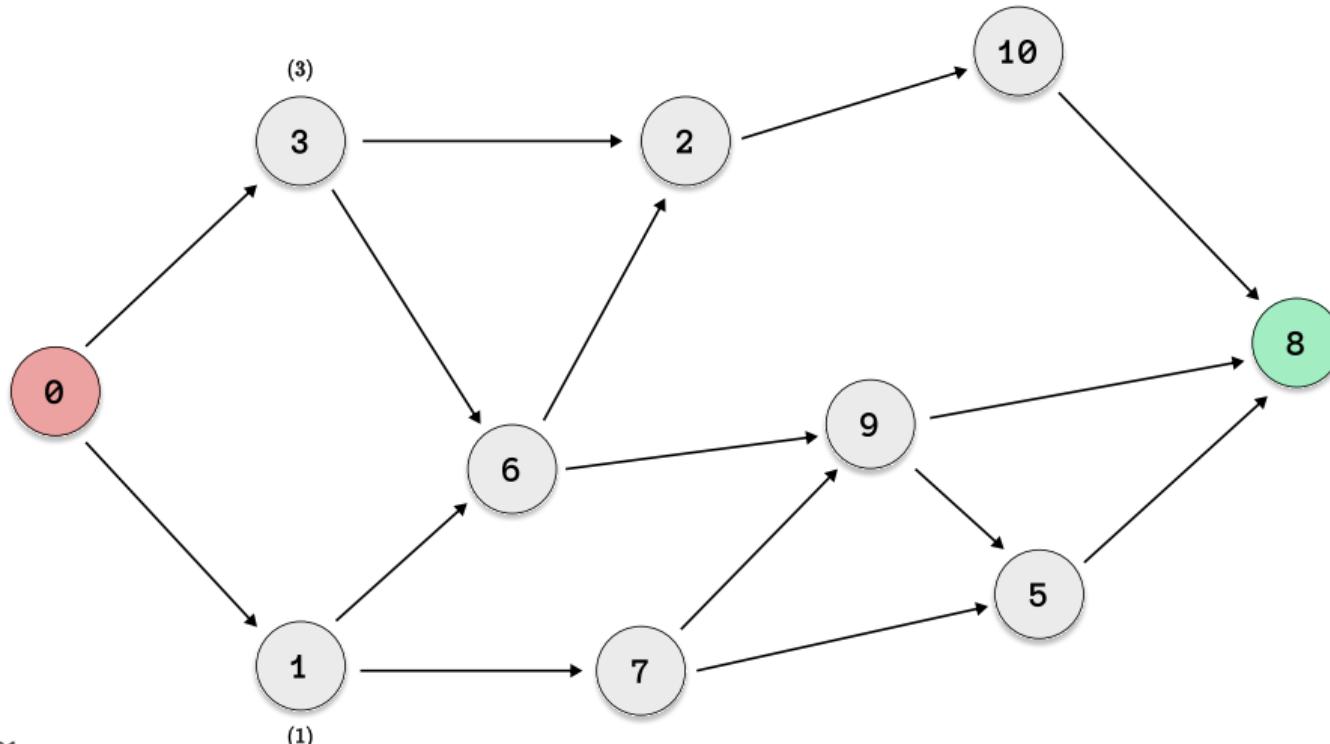
Visualizing the Recursive Definition





O-Set Construction Example

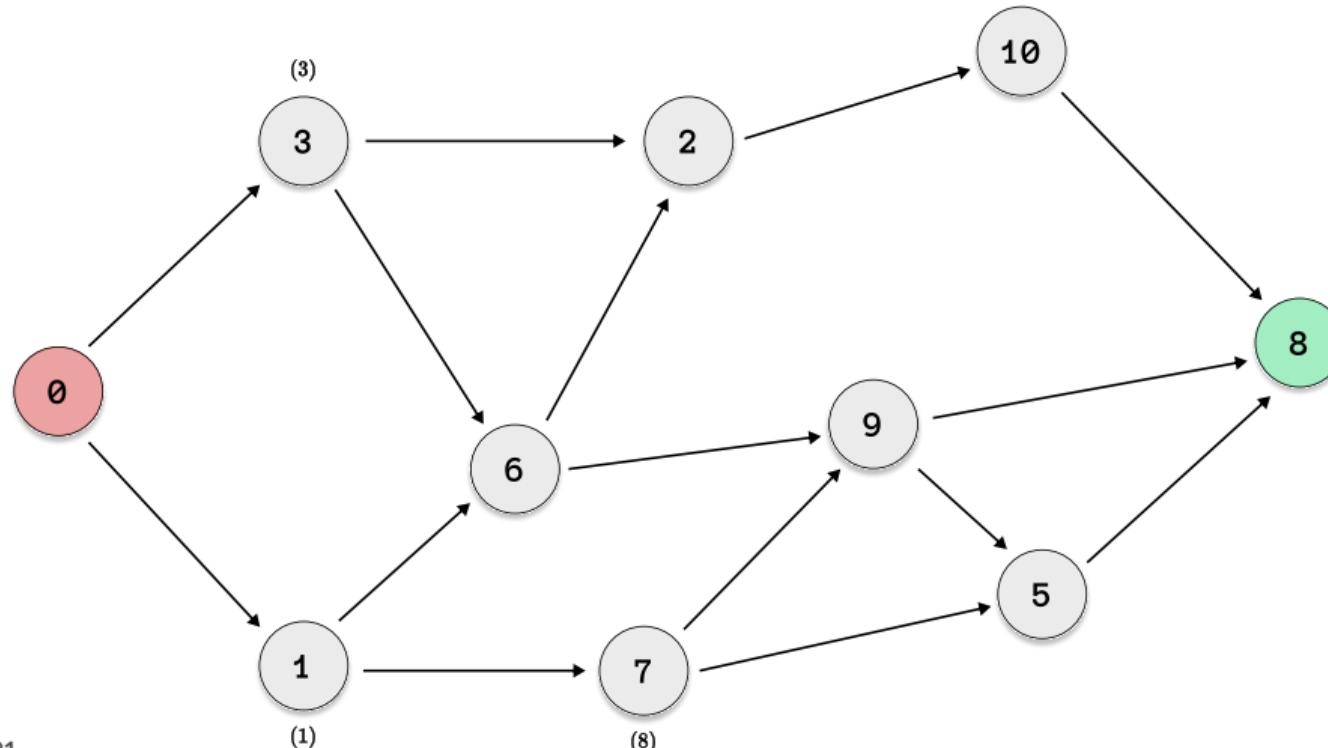
Visualizing the Recursive Definition





O-Set Construction Example

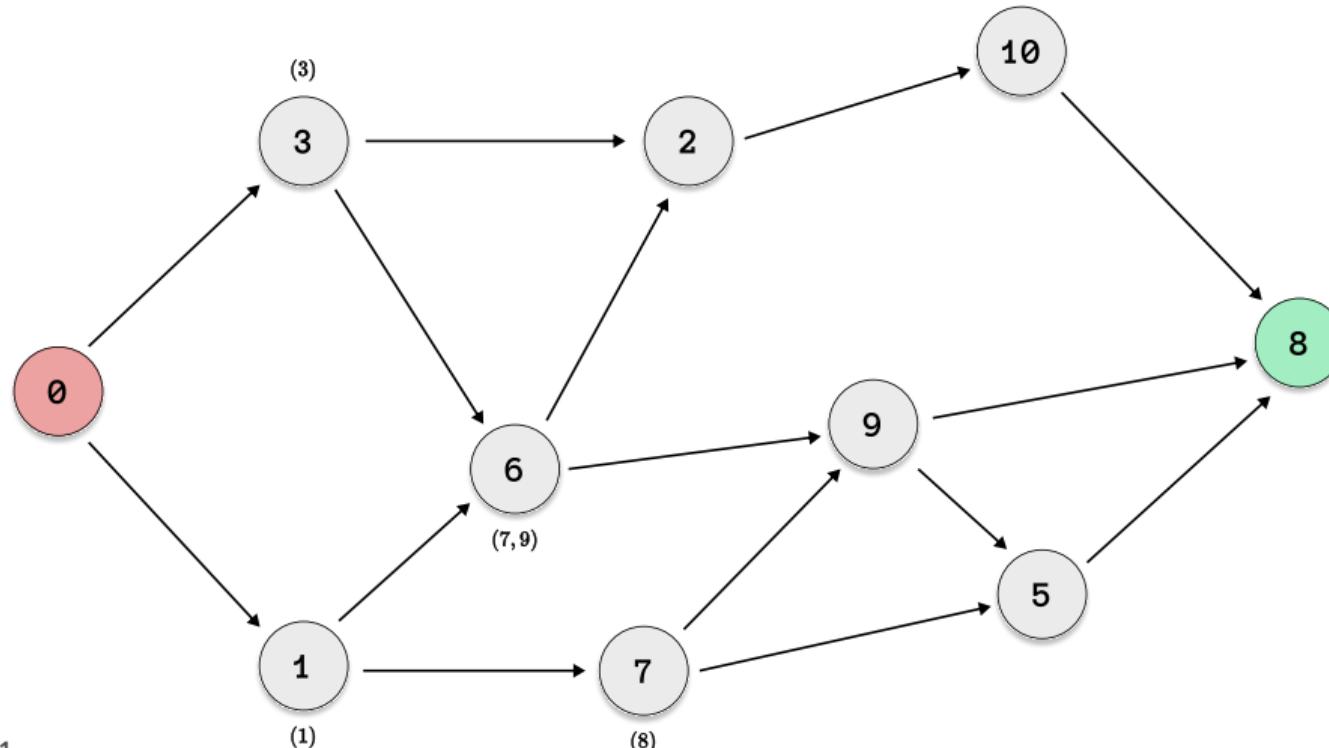
Visualizing the Recursive Definition





O-Set Construction Example

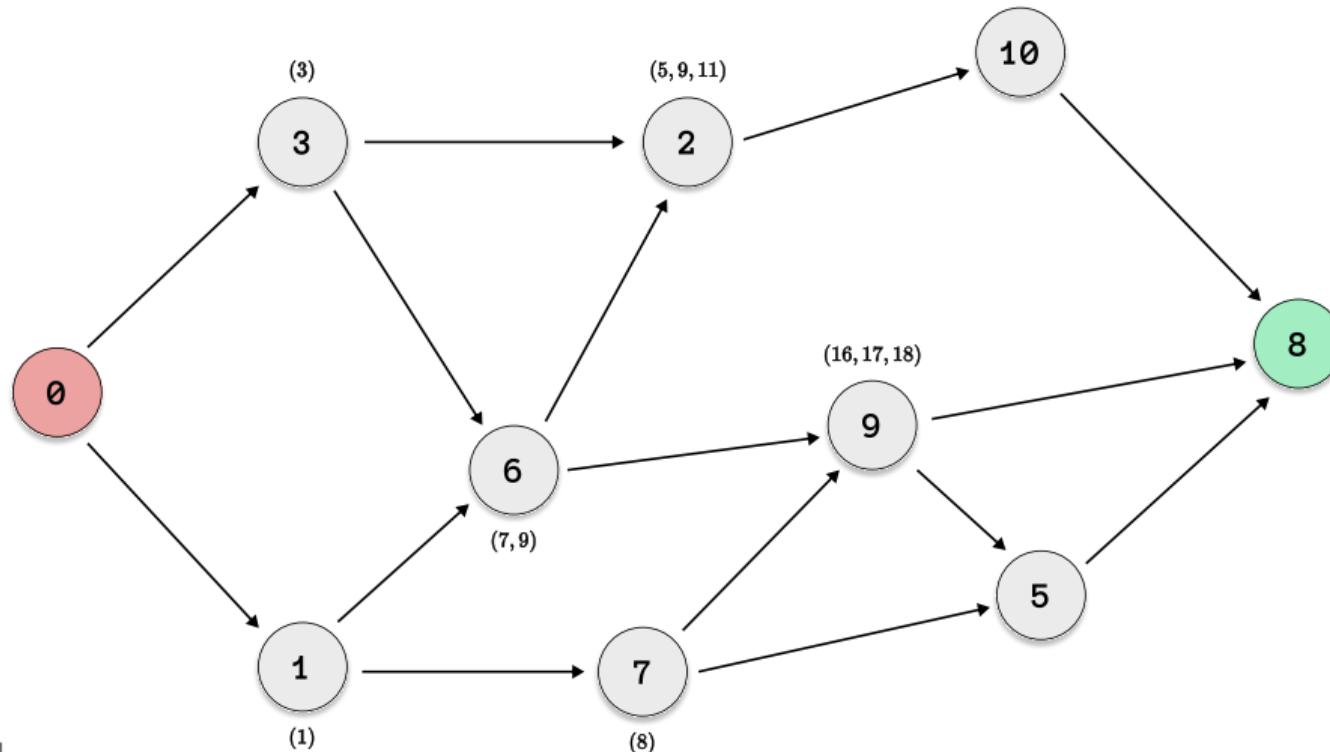
Visualizing the Recursive Definition





O-Set Construction Example

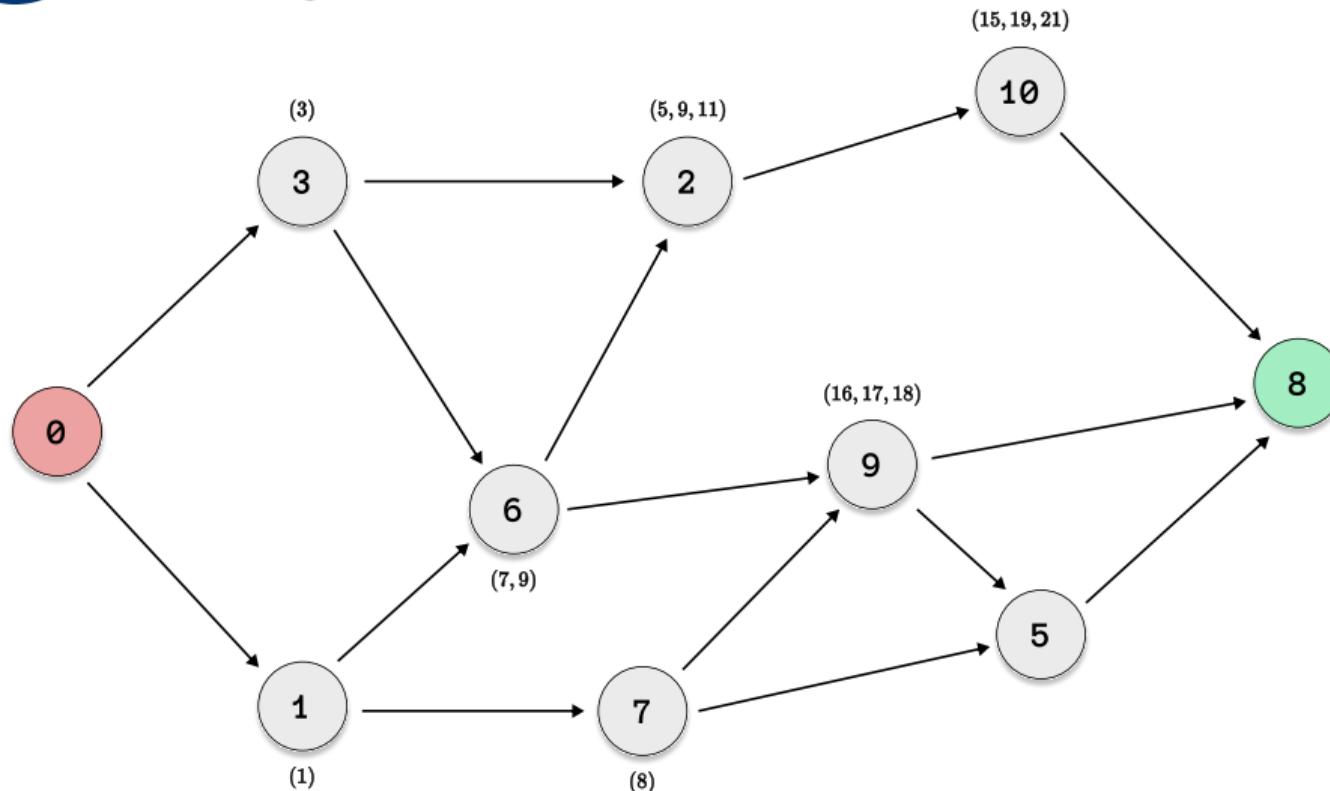
Visualizing the Recursive Definition





O-Set Construction Example

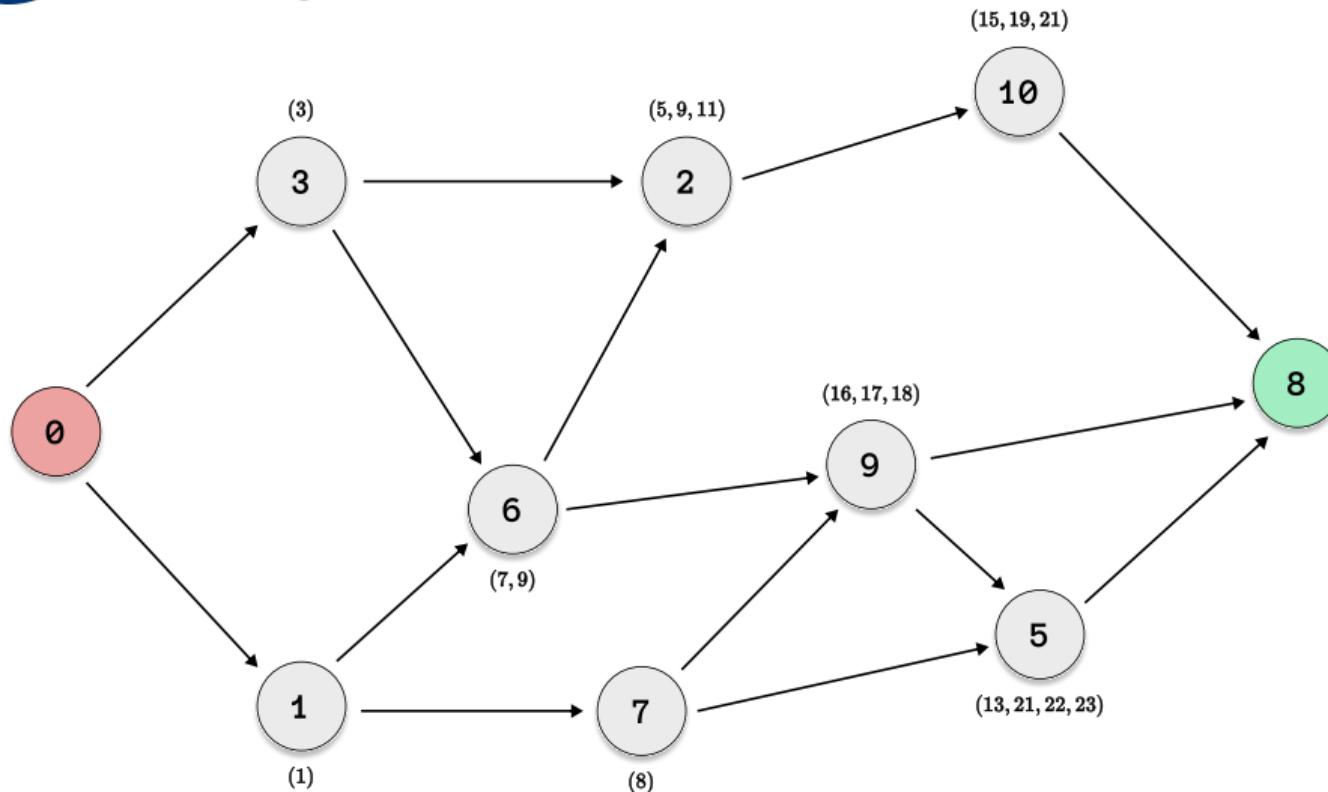
Visualizing the Recursive Definition





O-Set Construction Example

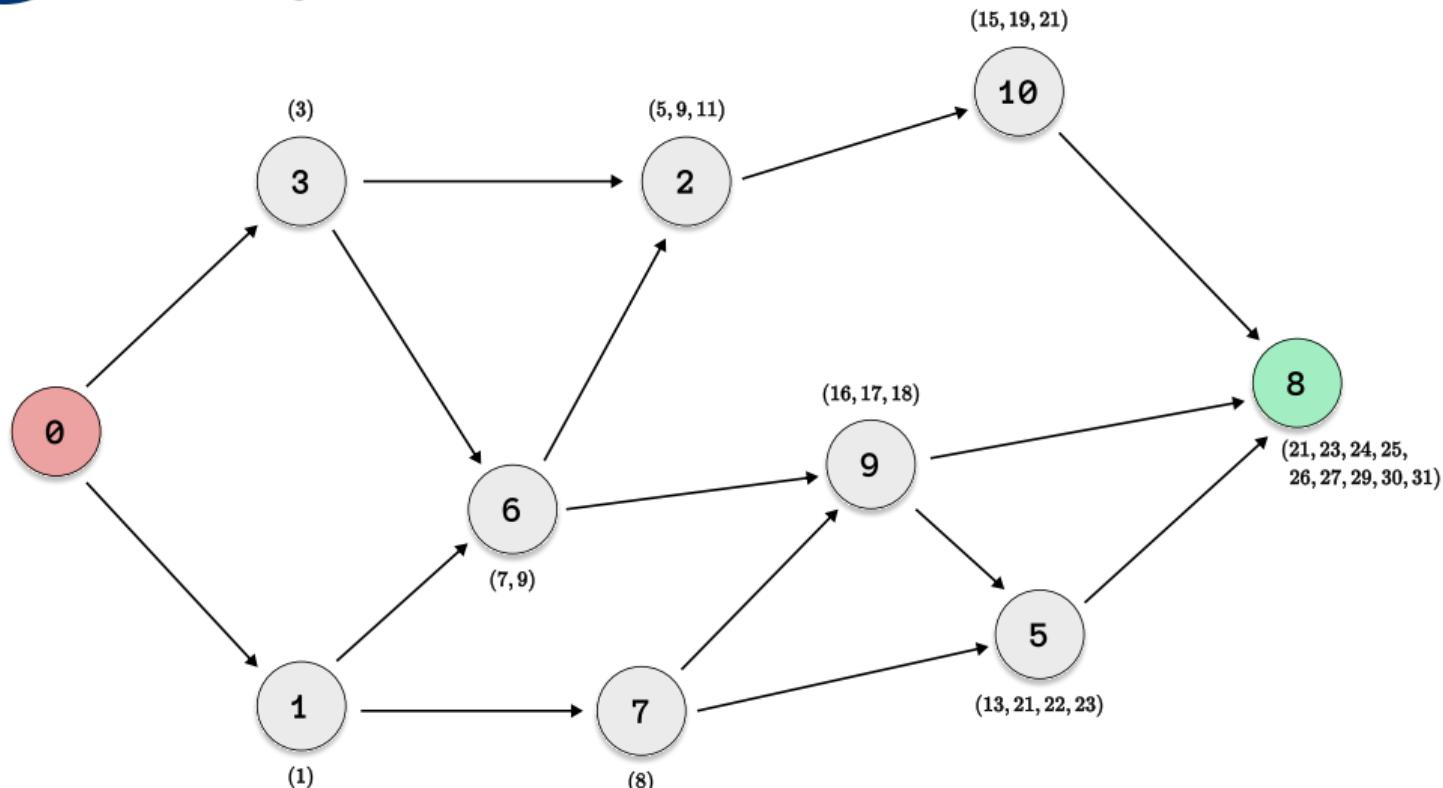
Visualizing the Recursive Definition





O-Set Construction Example

Visualizing the Recursive Definition





The Rank Query on Weighted DAGs

What Values are "Active" at Node N?

Rank Query on a Node N: $\text{rank}_{\mathcal{G}}(N)$

1. Returns a representation of a set of integers derived from the \mathcal{O} -set \mathcal{O}_N .

$$S_N = \bigcup_{x \in \mathcal{O}_N} \{z \in \mathbb{N}_0 \mid \max(0, x - w(N) + 1) \leq z \leq x\}.$$



The Rank Query on Weighted DAGs

What Values are "Active" at Node N?

Rank Query on a Node N: $\text{rank}_{\mathcal{G}}(N)$

1. Returns a representation of a set of integers derived from the \mathcal{O} -set \mathcal{O}_N .

$$S_N = \bigcup_{x \in \mathcal{O}_N} \{z \in \mathbb{N}_0 \mid \max(0, x - w(N) + 1) \leq z \leq x\}.$$

2. These intervals are then maximally merged. The query $\text{rank}_{\mathcal{G}}(N)$ returns a **minimal collection of disjoint closed integer intervals**

$$\mathcal{R}_N = \{[l_1, r_1], [l_2, r_2], \dots, [l_p, r_p]\}$$

such that their union exactly covers S_N .

\mathcal{R}_N captures the range of possible cumulative sums during the *activity* at node N



The Challenge: Storing Path Information

\mathcal{O} -Sets Can Be Huge!

- **Problem:** The size $|\mathcal{O}_v|$ can grow very large!
- **Question:** Can we represent the necessary information more compactly?



The Challenge: Storing Path Information

\mathcal{O} -Sets Can Be Huge!

- **Problem:** The size $|\mathcal{O}_v|$ can grow very large!
- **Question:** Can we represent the necessary information more compactly?

Core Idea: Partitioning + Indirection

Partition vertices V into two types:

1. Explicit Vertices (V_E)

Store \mathcal{O}_v directly.
(Simple, but potentially large)

2. Implicit Vertices (V_I)

Do not store \mathcal{O}_v explicitly
(Reconstruct on-the-fly.)



The Challenge: Storing Path Information

\mathcal{O} -Sets Can Be Huge!

- **Problem:** The size $|\mathcal{O}_v|$ can grow very large!
- **Question:** Can we represent the necessary information more compactly?

Core Idea: Partitioning + Indirection

Partition vertices V into two types:

1. Explicit Vertices (V_E)

Store \mathcal{O}_v directly.
(Simple, but potentially large)

2. Implicit Vertices (V_I)

Do not store \mathcal{O}_v explicitly
(Reconstruct on-the-fly.)

Reconstruction for $v \in V_I$ using:

- Designated Successor $\sigma(v)$
- Offset Sequence \mathcal{J}_v (at v)



Implicit Reconstruction: Successor & Offset

How V_I Nodes Refer to Others

1. Designated Successor $\sigma(v)$ (for $v \in V_I$)

Which successor should v point to? **Heuristic:** Choose $u = \sigma(v)$ that minimizes $|\mathcal{O}_u|$.

$$\sigma(v) \in \operatorname{argmin}_{u \in \text{Succ}(v)} \{|\mathcal{O}_u|\}.$$

2. Offset Sequence \mathcal{J}_v (for $v \in V_I$)

How to get \mathcal{O}_v from $\mathcal{O}_{\sigma(v)}$? Let $u = \sigma(v)$.

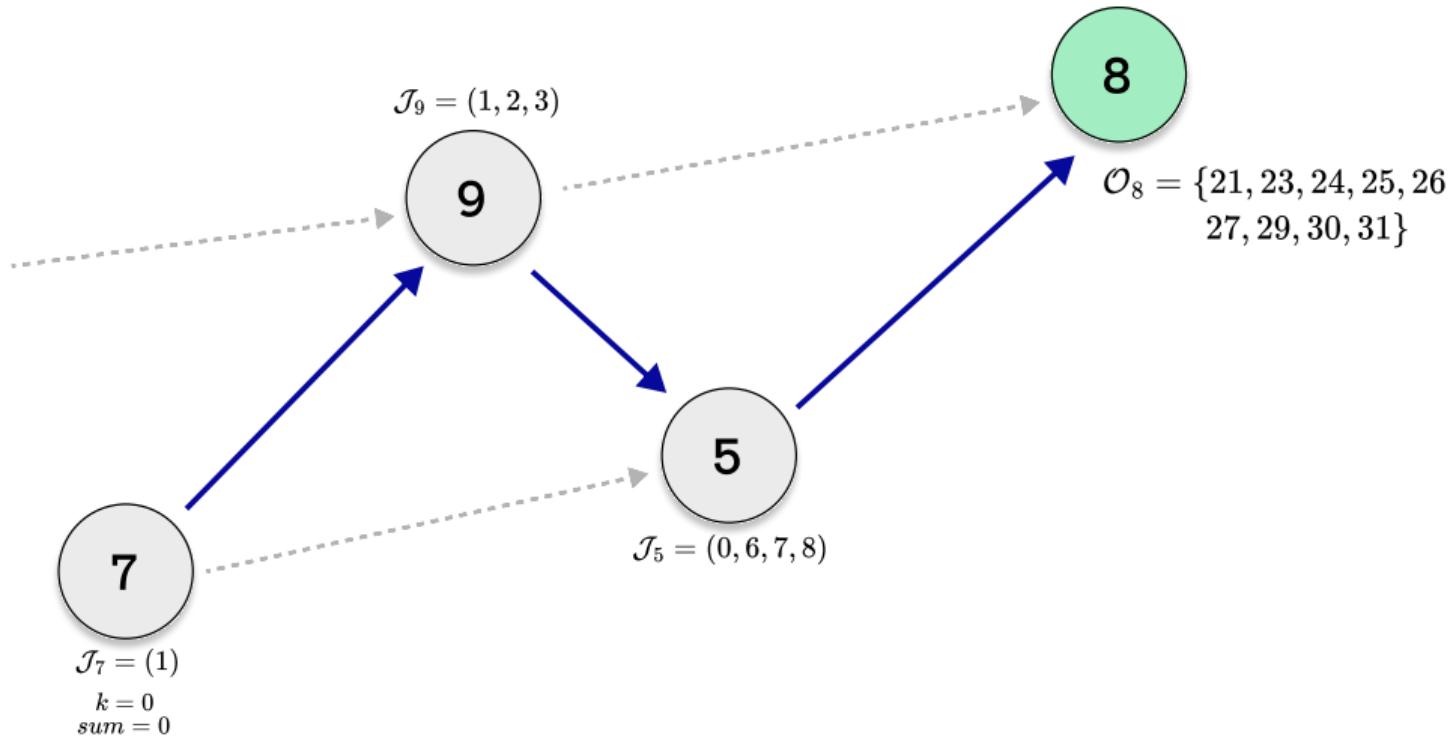
- **Relationship:** Each element $x_k \in \mathcal{O}_v$ comes from some $y_{j_k} \in \mathcal{O}_u$ via $x_k = y_{j_k} - w(u)$.
- **Offset Sequence \mathcal{J}_v :** Stores the index j_k corresponding to each x_k .

$$\mathcal{J}_v = (j_0, j_1, \dots, j_{m-1}), \quad \text{where } m = |\mathcal{O}_v|$$



Example: Computing $\mathcal{O}_7[0]$

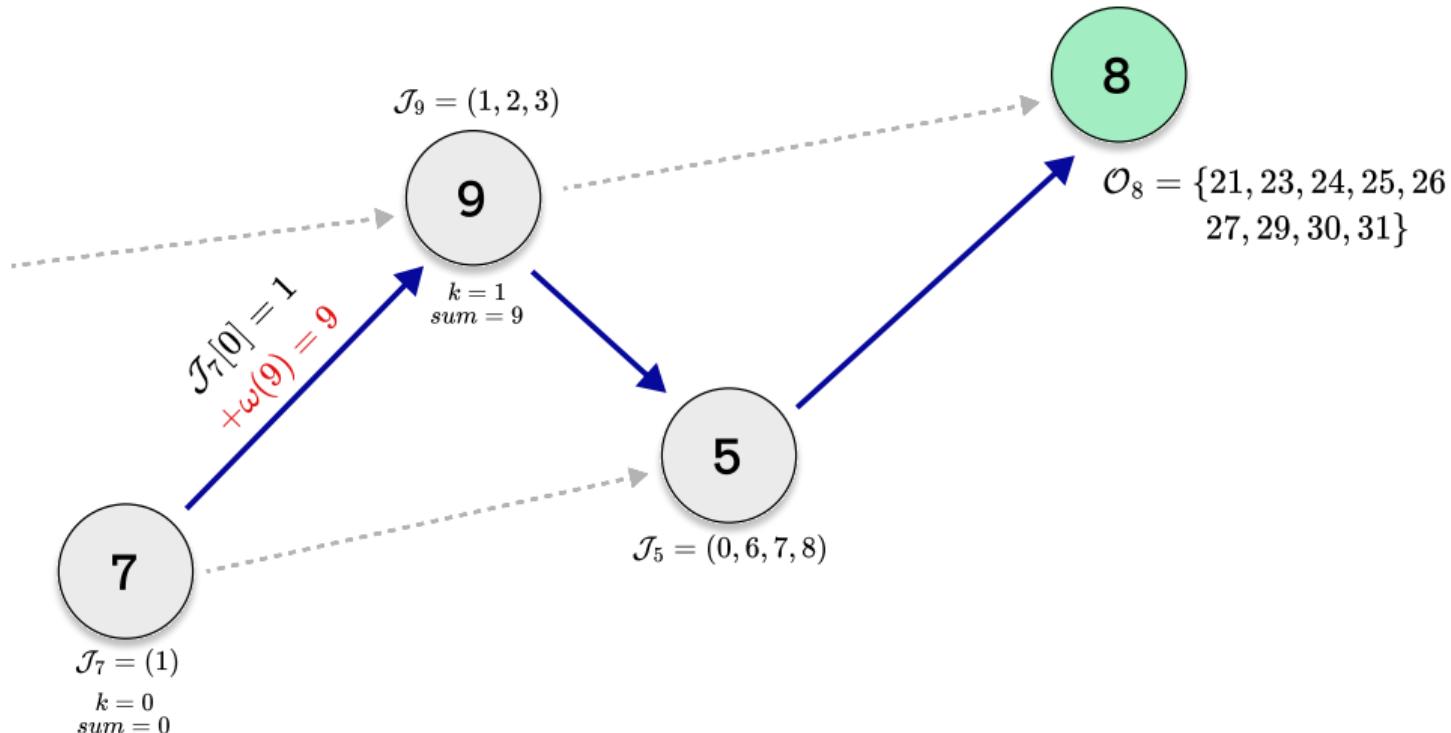
Following the Successor Path: $7 \rightarrow 9 \rightarrow 5 \rightarrow 8$





Example: Computing $\mathcal{O}_7[0]$

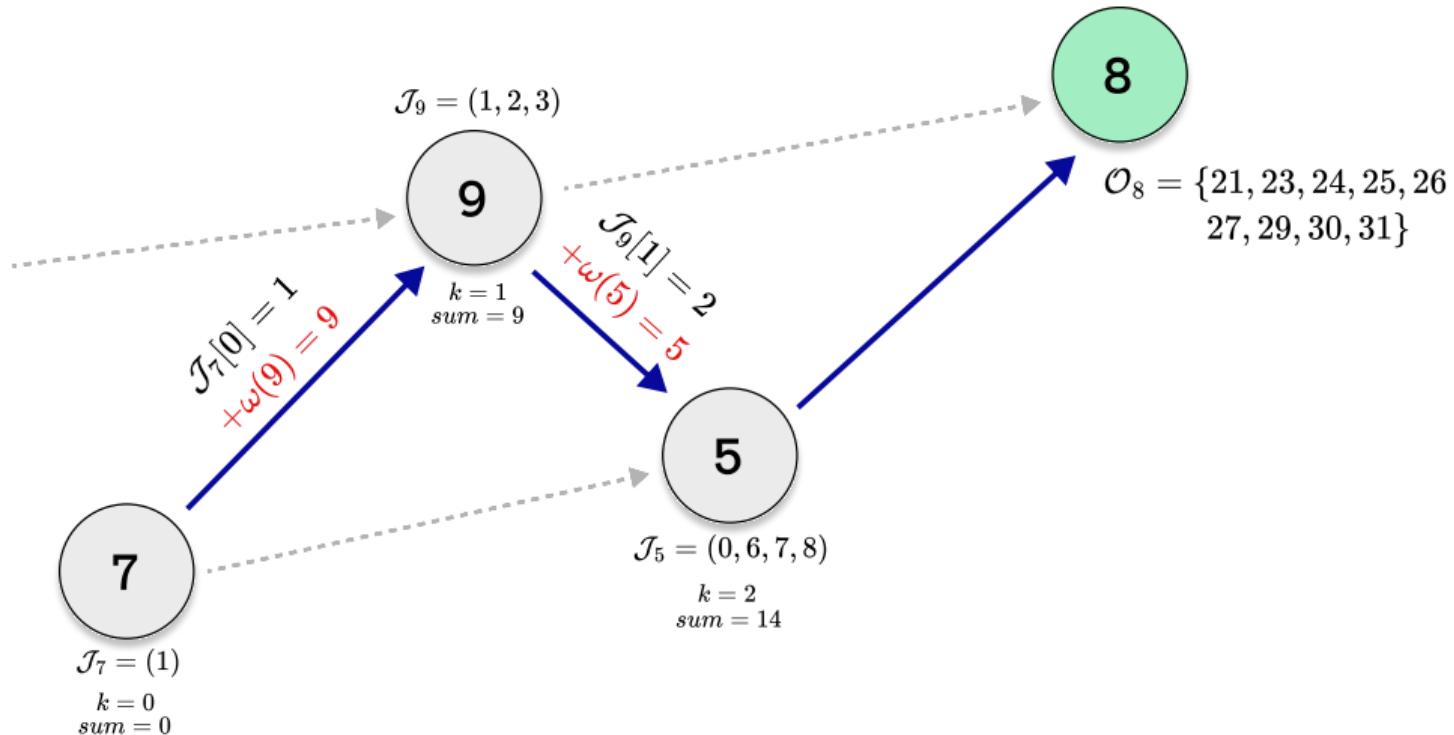
Following the Successor Path: $7 \rightarrow 9 \rightarrow 5 \rightarrow 8$





Example: Computing $\mathcal{O}_7[0]$

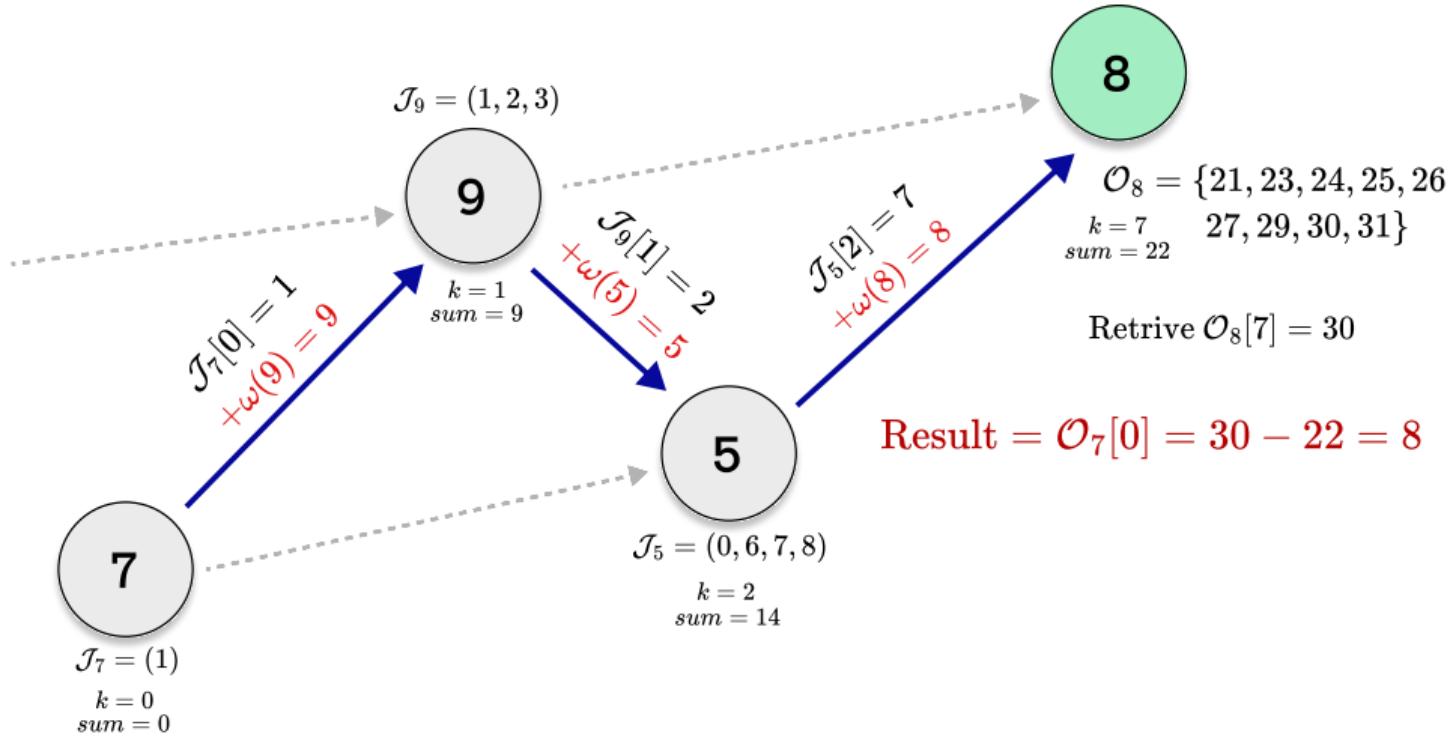
Following the Successor Path: $7 \rightarrow 9 \rightarrow 5 \rightarrow 8$





Example: Computing $\mathcal{O}_7[0]$

Following the Successor Path: $7 \rightarrow 9 \rightarrow 5 \rightarrow 8$

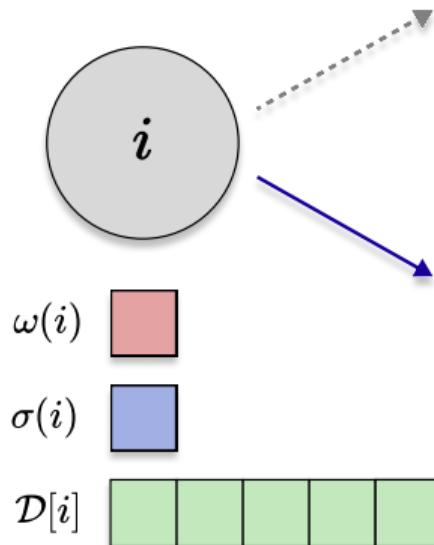




Succinct Data Structure: Components

Arrays Indexed by Vertex ID

Each node stores 3 components

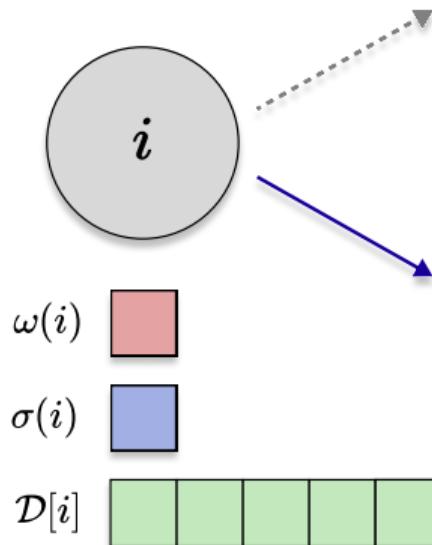




Succinct Data Structure: Components

Arrays Indexed by Vertex ID

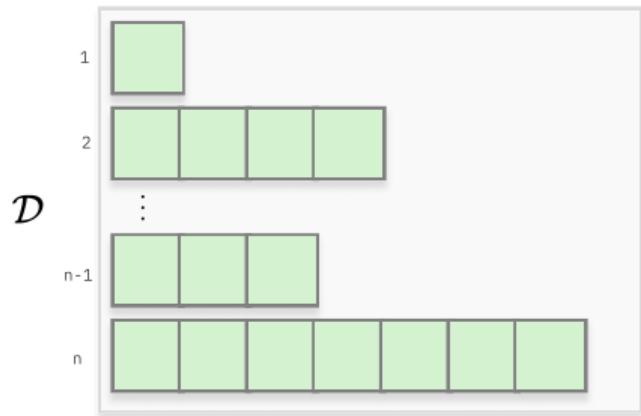
Each node stores 3 components



Succinct DAG as a Struct of Arrays

ω	0	$\omega(1)$	$\omega(2)$	\dots	\dots	\dots	\dots	$\omega(n)$
----------	---	-------------	-------------	---------	---------	---------	---------	-------------

σ	0	$\sigma(1)$	$\sigma(2)$	\dots	\dots	\dots	\dots	$\sigma(n)$
----------	---	-------------	-------------	---------	---------	---------	---------	-------------





Compression Strategies

Reducing Memory Footprint

Component	Description
\mathcal{W} (Node weights)	Array of positive integers.
Σ (Successor IDs)	Array of positive integers.



Compression Strategies

Reducing Memory Footprint

Component	Description
\mathcal{W} (Node weights)	Array of positive integers.
Σ (Successor IDs)	Array of positive integers.

Compression Options

- Variable-Length Integer Coding —→ we published a Rust library^a for this!
- Wavelet Trees (*for nodes with small weight range*)

^a<https://crates.io/crates/compressed-intvec>



Compression Strategies

Reducing Memory Footprint

Component	Description
\mathcal{W} (Node weights)	Array of positive integers.
Σ (Successor IDs)	Array of positive integers.
\mathcal{D} (Associated Data)	Array of arrays of increasing positive integers.

Compression Options

- Variable-Length Integer Coding —→ we published a Rust library^a for this!
- Wavelet Trees (*for nodes with small weight range*)

^a<https://crates.io/crates/compressed-intvec>



Compression Strategies

Reducing Memory Footprint

Component	Description
\mathcal{W} (Node weights)	Array of positive integers.
Σ (Successor IDs)	Array of positive integers.
\mathcal{D} (Associated Data)	Array of arrays of increasing positive integers.

Compression Options

- Variable-Length Integer Coding —→ we published a Rust library^a for this!
- Wavelet Trees (*for nodes with small weight range*)
- Elias-Fano Encoding (*for monotonic sequences*)
- Run-Length Encoding (RLE) (*for clustered monotonic sequences*)

^a<https://crates.io/crates/compressed-intvec>



Space Efficiency: Baseline Comparison

How Much Information is in the Graph?

To evaluate our structure's space, we need a baseline.

0th-Order Graph Entropy $H_0(G)$

A theoretical lower bound for storing the *entire* weighted DAG (V, E, w) losslessly.

$$H_0(G) = \underbrace{H_W(G)}_{\text{Cost for Weights}} + \underbrace{H_E(G)}_{\text{Cost for Topology (Edges)}}$$



Space Efficiency: Baseline Comparison

How Much Information is in the Graph?

To evaluate our structure's space, we need a baseline.

0th-Order Graph Entropy $H_0(G)$

A theoretical lower bound for storing the *entire* weighted DAG (V, E, w) losslessly.

$$H_0(G) = \underbrace{H_W(G)}_{\text{Cost for Weights}} + \underbrace{H_E(G)}_{\text{Cost for Topology (Edges)}}$$

- $H_W(G) \approx \sum_{v \in V} \log(w(v) + 1)$ bits (*Minimal binary encoding for weights*).



Space Efficiency: Baseline Comparison

How Much Information is in the Graph?

To evaluate our structure's space, we need a baseline.

0th-Order Graph Entropy $H_0(G)$

A theoretical lower bound for storing the *entire* weighted DAG (V, E, w) losslessly.

$$H_0(G) = \underbrace{H_W(G)}_{\text{Cost for Weights}} + \underbrace{H_E(G)}_{\text{Cost for Topology (Edges)}}$$

- $H_W(G) \approx \sum_{v \in V} \log(w(v) + 1)$ bits (*Minimal binary encoding for weights*).
- $H_E(G) \approx \log \binom{n(n-1)}{m}$ bits (*Cost to choose $m = |E|$ edges out of all possible $n(n - 1)$*).



Space Efficiency: Baseline Comparison

How Much Information is in the Graph?

To evaluate our structure's space, we need a baseline.

0th-Order Graph Entropy $H_0(G)$

A theoretical lower bound for storing the *entire* weighted DAG (V, E, w) losslessly.

$$H_0(G) = \underbrace{H_W(G)}_{\text{Cost for Weights}} + \underbrace{H_E(G)}_{\text{Cost for Topology (Edges)}}$$

- $H_W(G) \approx \sum_{v \in V} \log(w(v) + 1)$ bits (*Minimal binary encoding for weights*).
- $H_E(G) \approx \log \binom{n(n-1)}{m}$ bits (*Cost to choose $m = |E|$ edges out of all possible $n(n - 1)$*).

Any method saving the *full* graph structure needs at least $H_0(G)$ bits!



Space Comparison: Succinct Structure vs. Baselines

Bitcoin DAG Example ($n \approx 22k, m \approx 50k$)

Method	Estimated Bits
Theoretical Lower Bound	1,525,730
Weights $H_W(G)$	60,824
Topology $H_E(G)$	1,464,906

Precomputed Rank Queries:

- Explicit Binary Storage
- Elias-Fano Compressed

Our Succinct DAG

- Weights \mathcal{W}
- Successors Σ
- Assoc. Data \mathcal{D} (RLE)



Space Comparison: Succinct Structure vs. Baselines

Bitcoin DAG Example ($n \approx 22k, m \approx 50k$)

Method	Estimated Bits
Theoretical Lower Bound	1,525,730
Weights $H_W(G)$	60,824
Topology $H_E(G)$	1,464,906

Precomputed Rank Queries:

Explicit Binary Storage	4,854,533
Elias-Fano Compressed	2,211,849

Our Succinct DAG

- Weights \mathcal{W}
- Successors Σ
- Assoc. Data \mathcal{D} (RLE)



Space Comparison: Succinct Structure vs. Baselines

Bitcoin DAG Example ($n \approx 22k, m \approx 50k$)

Method	Estimated Bits
Theoretical Lower Bound	1,525,730
Weights $H_W(G)$	60,824
Topology $H_E(G)$	1,464,906
<hr/>	
Precomputed Rank Queries:	
Explicit Binary Storage	4,854,533
Elias-Fano Compressed	2,211,849
<hr/>	
Our Succinct DAG	602,808
Weights \mathcal{W}	60,824
Successors Σ	297,700
Assoc. Data \mathcal{D} (RLE)	244,284



Space Comparison: Succinct Structure vs. Baselines

Bitcoin DAG Example ($n \approx 22k, m \approx 50k$)

Method	Estimated Bits
Theoretical Lower Bound	1,525,730
Weights $H_W(G)$	60,824
Topology $H_E(G)$	1,464,906
<i>Precomputed Rank Queries:</i>	
Explicit Binary Storage	4,854,533
Elias-Fano Compressed	2,211,849
Our Succinct DAG	602,808
Weights \mathcal{W}	60,824
Successors Σ	297,700
Assoc. Data \mathcal{D} (RLE)	244,284

Achieving Sub-Entropy Space: How?

Our structure is **lossy** regarding the full graph topology:

- It **does not store** the complete edge set.
- It only stores the chosen successor $\sigma(v)$ for each implicit node (in Σ).

However, it is **lossless** for computing the specific **Rank Query**.



Future Direction: Bounded Query Time

Guaranteeing Predictable Performance

Performance Consideration

Query time for implicit node v depends on the length of the successor path

$$v \rightarrow \sigma(v) \rightarrow \sigma(\sigma(v)) \rightarrow \dots \rightarrow e \in V_E$$

Problem: Can be large/variable in deep DAGs \implies slow worst-case query time.



Future Direction: Bounded Query Time

Guaranteeing Predictable Performance

Performance Consideration

Query time for implicit node v depends on the length of the successor path

$$v \rightarrow \sigma(v) \rightarrow \sigma(\sigma(v)) \rightarrow \dots \rightarrow e \in V_E$$

Problem: Can be large/variable in deep DAGs \implies slow worst-case query time.

Solution, Challenges & Trade-offs

- **Solution:** Ensure every implicit node can reach an explicit node within k steps.
- **Challenges:** Finding the smallest possible V'_E that satisfies this condition is NP-hard (*minimum distance- k dominating set*).
- **Trade-off:** More explicit nodes \implies faster queries, but larger space.



Efficient Succinct Data Structures on Directed Acyclic Graphs

Thank you for listening!



Worst-Case \mathcal{O} -Set Size: Is Exponential Growth Possible?

Understanding the \mathcal{O} -set Size

Exponential Growth Can Occur

The cardinality of an \mathcal{O} -set, $|\mathcal{O}_v|$, is not generally bounded by a polynomial in the number of vertices $|V|$. It can grow exponentially.

Underlying Reason: Path Count

The number of distinct paths from a source s to a vertex v , denoted $|Path(s, v)|$, can itself be exponential in certain DAG structures. Since $|\mathcal{O}_v| \leq |Path(s, v)|$, the potential for exponential size exists.



Worst-Case \mathcal{O} -Set Size: Is Exponential Growth Possible?

Understanding the \mathcal{O} -set Size

Exponential Growth Can Occur

The cardinality of an \mathcal{O} -set, $|\mathcal{O}_v|$, is not generally bounded by a polynomial in the number of vertices $|V|$. It can grow exponentially.

Underlying Reason: Path Count

The number of distinct paths from a source s to a vertex v , denoted $|Path(s, v)|$, can itself be exponential in certain DAG structures. Since $|\mathcal{O}_v| \leq |Path(s, v)|$, the potential for exponential size exists.

Key Condition for Exponential Growth

The exponential potential is realized if the vertex weights $w(v)$ are assigned such that distinct paths $P_1 \neq P_2$ almost always lead to distinct cumulative weights $W(P_1) \neq W(P_2)$.



Achieving Exponential \mathcal{O} -Set Size

A Strategy for Path Weight Uniqueness

Start with a DAG structure that naturally admits an exponential number of paths between two nodes. An example is a layered graph with multiple choices at each layer transition.

Strategic Weight Assignment

Assign vertex weights $w(v)$ carefully to ensure path weight uniqueness.

$$w(v) = 2^k \quad (\text{using a unique exponent } k \text{ for each node})$$



Achieving Exponential \mathcal{O} -Set Size

A Strategy for Path Weight Uniqueness

Start with a DAG structure that naturally admits an exponential number of paths between two nodes. An example is a layered graph with multiple choices at each layer transition.

Strategic Weight Assignment

Assign vertex weights $w(v)$ carefully to ensure path weight uniqueness.

$$w(v) = 2^k \quad (\text{using a unique exponent } k \text{ for each node})$$

Mechanism: Unique Binary Representation

With power-of-2 weights, the cumulative path weight $W(P) = \sum_{v \in P \setminus \{s\}} w(v)$ becomes a sum of distinct powers of 2. Due to the uniqueness of binary representation, different sets of nodes (i.e., different paths) produce different sums. Therefore, $|Path(s, v)|$ distinct paths yield $|\mathcal{O}_v| = |Path(s, v)|$ distinct weights.